

Smart Parking Lot Assist System

Interim Report

Rajeev Atla, Parshva Mehta, Aman Patel, Abhiram Vemuri

Advisor: Kristin Dana

Team Number: SP25-02

I. INTRODUCTION

As transport infrastructure grows increasingly complex, large universities face mounting challenges in meeting the parking needs of students and faculty, especially with the continued reliance on individual car commuting. This project proposes an integrated solution using security camera imaging, computer vision, and a simple mobile application to ease the process of finding parking. The solution consists of two key components. First, a custom-configured security camera will be installed to overlook campus parking lots, capturing high-resolution images. These images are processed through a computer vision model, pre-trained to detect open parking spaces with high accuracy. The output from this model is then converted into a dynamic dataset identifying available parking spots in real-time. Second, this dataset is seamlessly integrated into a user-friendly mobile application focused on clarity and ease of use. The app features an adaptive interface with real-time maps, filters, and notifications, enabling users to quickly locate free parking spaces near their destinations and make more informed parking decisions. Past attempts to improve parking with automated valet systems have failed, often because of high costs, complicated logistics, and issues with reliability. In a lively campus setting, these valet systems could be more feasible: changing class schedules, various modes of transportation, and limited staff availability make it hard to provide reliable valet services. As a result, drivers waste precious time driving in circles in busy lots, leading to congestion, frustration, and adverse environmental effects. By combining advanced imaging technology with intuitive design, this project aims to enhance the daily commuting experience for students and faculty. Moreover, integrating this information into an accessible platform contributes to the broader goal of fostering sustainable and efficient urban mobility patterns, paving the way for data-driven transportation solutions on and beyond university campuses.

II. APPROACH METHODS

A. Dataset Prep

A diverse and robust dataset will form the foundation of this project. The dataset will include images of parking lot images of various occupancy levels. Images will be sourced from public datasets and proprietary captures using our camera to ensure relevance to the application. The preprocessing phase will involve resizing, normalization, and augmentation

to prepare the data for model training and ensure data quality and robustness. These processes help simulate real-world scenarios and address potential image quality variations caused by environmental factors. Annotation tools will meticulously label parking spaces as occupied or vacant. Specific attention will be given to special categories like handicapped spots, electric vehicle charging spaces, and reserved parking areas, ensuring the system's utility across diverse user needs. These steps will ensure the datasets' quality and reliability for training, setting the stage for robust model performance.

B. ML Model

[4] <https://docs.ultralytics.com/models/yolo11/>

The project will leverage the YOLOv11 computer vision algorithm [4] to address the task of parking space detection. Known for its accuracy and speed, YOLO is well-suited for real-time applications. The model will undergo customization and fine-tuning to align with the unique challenges of parking detection, such as variations in car sizes, overlapping boundaries, and partial occlusions caused by objects or shadows. By implementing advanced image processing techniques, the system will accurately identify parking spaces and determine whether they are vacant or occupied, even under diverse weather conditions and potential image quality disruptions. Training will focus on optimizing detection accuracy while minimizing computational demands to ensure seamless integration into low-power devices like the Raspberry Pi. Techniques like quantization and pruning will reduce model complexity without compromising performance. Rigorous testing and validation will ensure the model's robustness across different environments, meeting the necessary standards for real-time operation. As previously mentioned, the datasets gathered in the data acquisition phase will be used to evaluate and refine the machine learning algorithm, improving its overall functionality and accuracy.

C. Hardware

Selecting a suitable Raspberry Pi and camera module will be critical to the project's success. We need hardware that maintains stability (no outages/connection issues) and quality so the model runs smoothly. We will have a mounting system to maintain minimal movement so the resulting images are consistent. The testing phase will evaluate multiple aspects of the system, including stability and quality, and model detection accuracy across diverse parking lot layouts. Extensive field

testing will help refine the integration and optimize the system for practical deployment. Suppose the Raspberry Pi system cannot run the machine learning model. In that case, we will pivot to running the model on an external computer and sending images from the camera to the computer. In this case, the Raspberry Pi will be used to connect via Wi-Fi or ethernet.

D. Software

The software component will integrate advanced image processing, machine learning, and user-facing functionalities. Captured images will undergo preprocessing steps such as noise reduction, color correction, and enhancement to ensure the highest possible input quality for the trained model. The inference pipeline will leverage the trained model to identify and classify real-time parking spaces, offering high accuracy and speed. A robust data transmission protocol, such as Wi-Fi or cellular communication, will facilitate seamless information sharing with a central server or directly with a mobile application. The mobile application will serve as the primary interface for end-users, providing real-time parking availability information through interactive maps, search functionality, and real-time notifications. Empty parking spaces will be highlighted on the interactive map for a given parking lot so that the user will receive this information. Using the machine learning algorithm and the data gathered from the camera, we can feed this into the backend of the mobile app function and apply a front-end user interface that connects with this information. Emphasis will be placed on designing a user-friendly experience that accommodates various devices and user preferences, ensuring wide adoption and usability.

III. CHALLENGES

Throughout the development of our project, we encountered several significant challenges related to regulation and the supply chain, which ultimately limited the performance and scope of our system. Initially, the project was designed to utilize an aerial camera perspective, similar to the dataset images (see Figure 4), in order to accurately capture parking lines and define bounding boxes around each space. However, due to regulatory constraints from Rutgers University and the State of New Jersey, we were not permitted to install cameras on elevated platforms or drones as originally planned. As a result, we had to revise our approach and instead placed a stationary camera on the 5th floor of the CoRE building. This new vantage point introduced several visibility issues, as many parking spots were blocked by vehicles already parked in front of them, reducing the clarity and reliability of our input images.

In addition to these regulatory limitations, we also faced delays in receiving essential hardware components. The camera's power supply and SD card arrived late in the semester, preventing us from fully setting up and testing the system on-site. This significantly limited our ability to collect real-world data from the intended parking lot, which in turn negatively impacted the accuracy and effectiveness of our trained model. The power supply and SD card arriving late has significantly

impacted the ability to continue building the system pipeline. The goal was to have the camera take a continuous video, house the YOLOv11 model on the Raspberry Pi itself for local computing, take images at regular intervals with bounding boxes overlaid, send images to a database, pull the information from these images, and update a website with this information. Our development timeline has been dramatically affected by these issues.

Particularly on the software end of the project, the delay in receiving the hardware components needed to execute our project has drastically changed the way our database is being updated. The output of our ML model program needs to be used to update the database, but due to not having access to the power supply and SD card to fully test the system, we had to develop a workaround method for simply getting real data from the parking lot. The workaround we resorted to was using our phone cameras, which resulted in some limitations compared to what was originally envisioned such as lower quality resolution, and inconsistent frame rates of mobile devices. This made it harder for the ML algorithm to detect vacant or occupied parking spaces, which can result in missed or incorrect detections affecting the output file of the ML algorithm after running it and using the updated CSV file to update the database. These inaccuracies propagated through the pipeline resulting in a few false positives, resulting in a few incorrect entries into the database.

Another problem that was encountered was when trying to optimize the refresh rate for database stability and accuracy. Initially our system was trying to update the database every 5 to 10 seconds from reading the CSV file, but this had several implications such as the database locking, writing conflicts, and also occasional crashes when running the server for frontend requests. In order to solve this the refresh interval was increased to 30 seconds to avoid instability in the database system. This eased the load on the database, kept the access conflicts with the frontend requests to a minimal amount, and allowed the ML model more time to process images and get higher quality outputs closer to the real-time scenario. Overall, this more refined approach allowed for a more sustainable balance between the backend functionality of the system and the real-time data generated by the ML model.

IV. RESULTS

During the model training phase, a fine-tuning strategy was employed due to the absence of prior knowledge regarding the optimal number of training epochs. This approach involved training the model for a fixed number of epochs to obtain an initial set of weights. These weights were then used to initialize the model for subsequent training cycles, effectively serving as a foundation for further refinement. This iterative weight generation and retraining process was repeated across multiple rounds, allowing for gradual performance improvements while mitigating the risk of overfitting or underfitting.

We begin by analyzing the training and validation loss curves to evaluate the learning behavior of the model throughout the training process. These plots provide insight into the

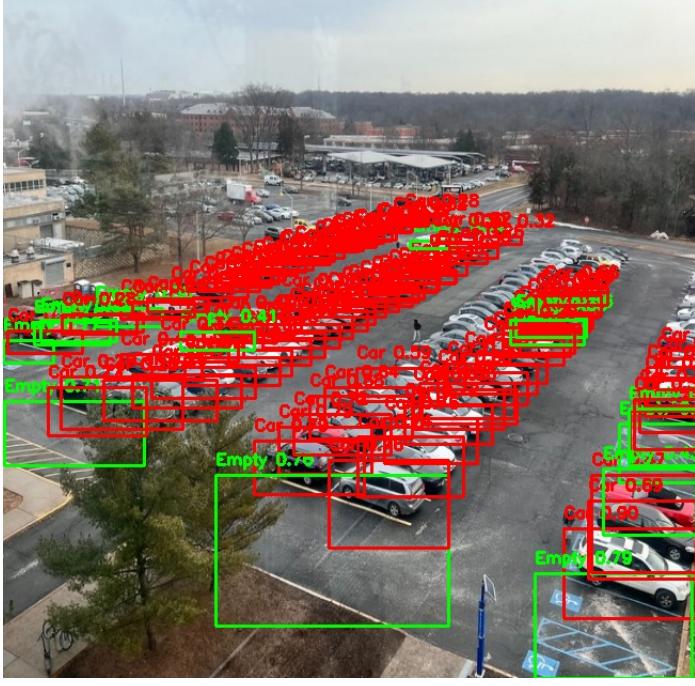


Fig. 5. Observe that the model struggles to identify vehicles and spaces at a distance. Most lots will need more than one camera in order to cover every space. This is a frame taken from a test video.

system's energy-efficient design also contributes to lower long-term operational costs, and potential tax incentives for smart city infrastructure, carbon footprint reduction, and energy-efficient technologies may reduce implementation expenses. Additionally, the system may provide indirect cost savings by lowering vehicle idling time, reducing fuel consumption, and minimizing emissions due to inefficient parking behaviors.

Environmentally, the system contributes positively by reducing greenhouse gas emissions associated with drivers circling to find parking. The hardware components, such as the Raspberry Pi, have a low power footprint and can operate efficiently with minimal energy consumption. Furthermore, the use of commonly available and recyclable materials supports a sustainable product lifecycle. As the system minimizes unnecessary vehicle movement, it indirectly reduces urban traffic congestion and resource overuse.

Socially, the impact of the project is notable in its ability to improve urban mobility, reduce driver frustration, and promote smarter city planning. By streamlining the parking process, it increases accessibility for all, including the elderly and those with disabilities, and reduces stress in high-traffic zones. While partial automation may reduce the need for manual lot attendants, it simultaneously opens avenues for new tech-related roles in installation, maintenance, and data analysis. The project also enhances public safety by reducing the chances of illegal or hazardous parking and complies with evolving regulations around smart infrastructure and urban sustainability.

VI. COMPUTER VISION/SOCIETY

- [1] <https://ieeexplore.ieee.org/document/10009811>
- [2] <https://ieeexplore.ieee.org/document/10914966>
- [3] <https://ieeexplore.ieee.org/document/10332925>

The spread of machine vision technologies has revolutionized urban infrastructure, positioning robotics at the forefront of societal innovation. These vision-based systems are integral to smart cities, driving automation in critical areas such as traffic management, security, and public services. By enabling real-time visual interpretation, these systems empower autonomous agents that improve urban life through enhanced efficiency, sustainability, and safety.

Recent advancements have shown that efficient, low-resource methods can now achieve vision-based object detection and tracking. An object-tracking approach utilizing frame-differencing, adaptive thresholding, and tracking algorithms, which enables real-time detection of moving objects with minimal computational overhead, was demonstrated [1]. Such techniques offer an affordable alternative to traditional hardware-intensive solutions, making intelligent automation more accessible to a broader range of applications. Concurrently, deep learning-based object detection systems, such as YOLO, have significantly advanced the capabilities of machine vision by balancing inference accuracy with computational efficiency, allowing real-time detection in everyday environments.

In the context of urban mobility, these advancements have been especially impactful in the development of smart parking systems. The work in [2] illustrates how machine vision can optimize parking space identification and classification, alleviating traffic congestion and maximizing space utilization in urban centers. Such systems can provide immediate, data-driven insights, improving user experience and urban planning. Automating routine tasks like parking identification reduces reliance on human labor and contributes to better traffic flow, environmental sustainability, and a more responsive urban environment.

Nevertheless, the widespread deployment of machine vision systems raises critical ethical concerns. Vision-based surveillance technologies, particularly in public spaces, present privacy, consent, and data security challenges. Unlike traditional surveillance systems, which often rely on biometric data, modern implementations like YOLO-based systems can focus solely on objects rather than individuals, offering a potential solution to privacy concerns. Despite this, future deployments must ensure privacy by design and transparency in how data is processed and used to maintain public trust and accountability.

Beyond urban mobility, integrating robotics into broader societal functions reflects a paradigm shift in public services. As [3] highlights, advances in social robotics show how robotic systems can enhance educational and medical services, particularly in pediatric and healthcare settings. Robots can assist in personalized learning, therapeutic activities, and medical monitoring by applying similar vision-based technologies. These advancements highlight the growing potential of

robotics to serve not just as tools for automation but as agents of positive social change, addressing gaps in critical areas like education and healthcare while reinforcing the need for ethical considerations in their design and deployment.

The integration of robotic perception into civic infrastructure signals a shift in the relationship between technology and society. These systems highlight the growing role of automation in providing public services, raising important questions about responsibility, inclusivity, and governance in their design and deployment. As machine vision and robotics continue to evolve, their potential to reshape urban environments will depend on technical capabilities and their alignment with societal values and priorities.

In this evolving landscape, even low-cost systems like YOLO-based parking detection provide valuable insights into how vision-enabled robotics can be deployed for the public good. By examining their practical limitations and benefits, such systems offer important lessons in integrating intelligent technologies into public infrastructure while addressing ethical concerns and improving urban living conditions.

VII. CONCLUSION

The project envisions a highly scalable and adaptable system capable of being deployed in a wide range of environments, including shopping malls, stadiums, airports, university campuses, and densely populated urban centers. Its modular design ensures seamless integration with existing infrastructure, making it suitable for both small-scale and large-scale parking facilities. By leveraging the portability and affordability of edge devices such as Raspberry Pi, the system remains cost-effective while maintaining reliable performance. This flexibility opens opportunities for widespread adoption, especially in developing regions or in contexts where large-scale infrastructure investments are not feasible.

Looking ahead, several enhancements are planned to increase the system's intelligence, responsiveness, and utility. These include integration with smart city frameworks to deliver real-time traffic and parking analytics, aiding not only drivers but also city planners and municipal authorities. Predictive modeling techniques using machine learning can be applied to forecast parking space availability based on historical usage patterns and live sensor inputs, optimizing traffic flow and minimizing time spent searching for parking. The architecture can also be extended into multi-camera or drone-assisted configurations, enabling coverage of vast and dynamic environments such as open-air lots or event venues. Drones on predetermined flight paths could provide aerial perspectives, increasing monitoring flexibility and responsiveness. By combining state-of-the-art machine vision with user-centric design principles, this system presents a forward-thinking, practical solution to the persistent challenges of urban parking, paving the way for smarter, more efficient cities.

VIII. APPENDIX - CODE